

$$\nabla \cdot \mathbf{D} = \rho$$

$$\nabla \cdot \mathbf{B} = 0$$

$$\nabla \times \mathbf{E} = -\frac{\partial \mathbf{B}}{\partial t}$$

$$\nabla \times \mathbf{H} = \mathbf{J} + \frac{\partial \mathbf{D}}{\partial t}$$

GPU Programming in Computational Electromagnetics

Yunlong Lian
@algorithmx

Outline

- Computational electromagnetics algorithms
 - Finite-Difference Time-Domain (FDTD)
 - Rigorous Coupled-Wave Analysis (RCWA)
 - Finite-Element Method (FEM)
- Optimization of package **RigorousCoupledWaveAnalysis.jl**



Computational electromagnetics algorithms

Computational electromagnetics algorithms

- Classification
 - Space
 - *differential* equations
 - *integral* equations
 - Time
 - time domain discretization
 - frequency domain

Computational electromagnetics algorithms

- Numerical solution of the Maxwell *differential* equations
 - Time domain
 - Uniform spatial discretization
 - Finite-Difference Time-Domain (FDTD)
 - Adaptive spatial discretization
 - Time-Domain Finite-Element Method (TDFEM)

Computational electromagnetics algorithms

- Numerical solution of the Maxwell *differential* equations
 - Frequency domain
 - Static (zero frequency)
 - Finite-Element Method (FEM)
 - Time-harmonic (single frequency)
 - Rigorous Coupled-Wave Analysis (RCWA)

Computational electromagnetics algorithms

- Numerical solution of the Maxwell *integral* equations
 - Boundary-Element Method (BEM)
 - Method of Moments (MoM)
 - ...

Finite-Element Method (FEM)

- Noticeable open-source projects
 - Netgen/NGSolve : <https://ngsolve.org/>
 - FEniCSx : <https://fenicsproject.org/>
 - libMesh : <http://libmesh.github.io/>
 - FreeFEM : <https://freefem.org/>
- Many commercial softwares
 - COMSOL
 - ANSYS

Finite-Element Method (FEM)

- Basic idea: approximate field in continuum by values on finite elements (via Galerkin method)
 - Field \rightarrow vector
 - Differential operator \rightarrow linear operator
 - Boundary condition \rightarrow constraints
- Special considerations for electromagnetic field and Maxwell equations
 - Vector field

Finite-Element Method (FEM)


- Technical challenges and potential GPU accelerations
 - Mesh generation
 - Mesh quality control
 - Sparse solver

Finite-Difference Time-Domain (FDTD)

- Noticeable open-source projects
 - gprMax : <https://github.com/gprMax/gprMax>
 - mumax3 : <http://mumax.github.io/>
 - Python 3D FDTD Simulator : <https://github.com/flaport/fdtd>
 - Meep : <https://github.com/NanoComp/meep>
 - Tidy3D (commercial, FlexCompute Inc.) :
<https://github.com/flexcompute/tidy3d>

Finite-Difference Time-Domain (FDTD)

- Straightforward discretization of time dependent Maxwell equation
- Yee lattice



Finite-Difference Time-Domain (FDTD)

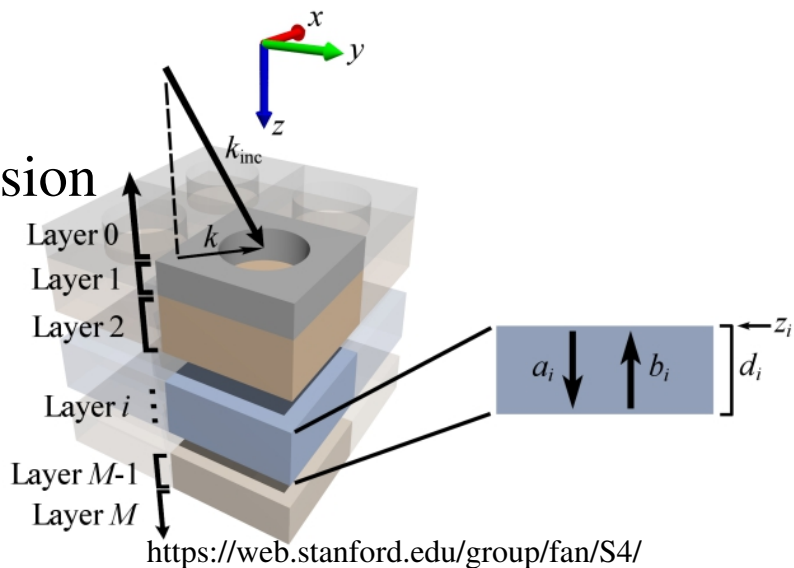
- Technical challenges and potential GPU accelerations
 - Memory cost cannot be reduced
 - Memory access patterns in updating the Yee lattice
 - Grid communications for parallelized algorithm

Rigorous Coupled-Wave Analysis (RCWA)

- Noticeable open-source projects
 - S4, or Stanford Stratified Structure Solver :
<https://github.com/victorliu/S4>
 - GRCWA : <https://github.com/weiliangjinca/grcwa>
 - EMPossible course : <https://github.com/zhaonat/Rigorous-Coupled-Wave-Analysis>
 - Jordan Edmunds : <https://github.com/edmundsj/RCWA>
 - RigorousCoupledWaveAnalysis.jl :
<https://github.com/jonschlipf/RigorousCoupledWaveAnalysis.jl>

Rigorous Coupled-Wave Analysis (RCWA)

- Enhanced Transmission Matrix algorithm by Moharam
 - major steps
 - partition the film into layers
 - perform layerwise 2D Fourier transform
 - calculate the wave amplitudes
 - along the propagation, for transmission
 - backwards, for reflection
 - collect results



Rigorous Coupled-Wave Analysis (RCWA)

- Technical challenges and potential GPU accelerations
 - Accuracy vs efficiency: N = FT orders, $O(N^6)$ complexity, eigen solver $O(M^3)$ matrix dimension $O(N^2)$ by $O(N^2)$
 - Conformity in description of shapes
 - Multiple films on curved surface



Experiment: RCWA

Experiment: RCWA

- Motivation
 - Wide range of applications
 - Nanophotonics (academic)
 - Optical Critical Dimension (OCD) for process control in semiconductor manufacturing (industrial)
 - RCWA has to be fast
- **Use Julia and CUDA.jl**
- Starting point
 - <https://github.com/jonschlipf/RigorousCoupledWaveAnalysis.jl>
 - Rigorous coupled-wave analysis of a multi-layered plasmonic integrated refractive index sensor, Opt. Express 29, 36201-36210 (2021)
- Project GitHub repository

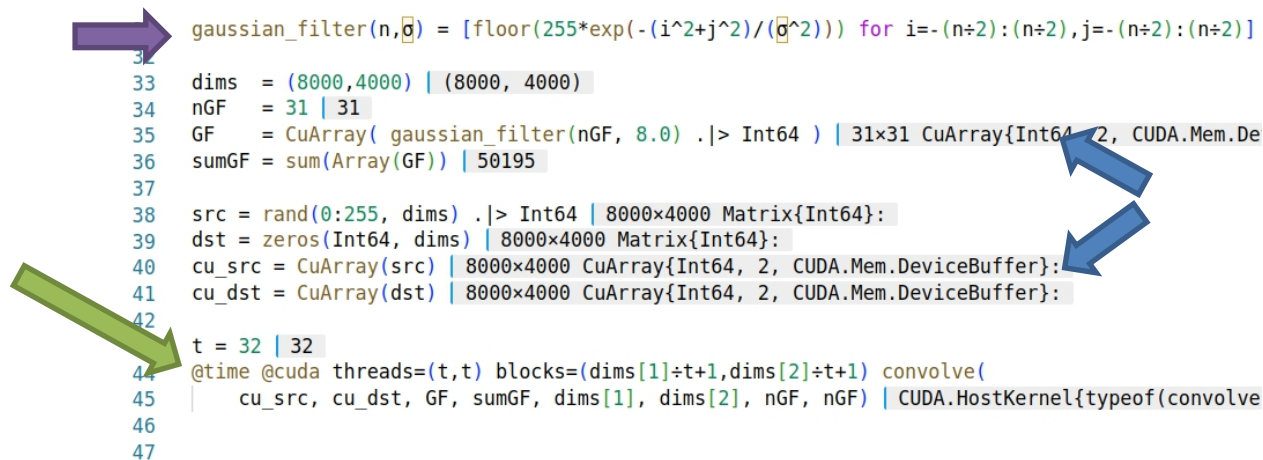
Why coding with Julia and CUDA.jl?

- Features of Julia programming language
 - garbage collection
 - array 1-based
 - support multiple dispatch and meta-programming
 - convenient profiling (@profview + VSCode)
- Features of CUDA.jl
 - using Julia syntax and grammar to **write, compile and run** CUDA kernels
 - easy memory management
 - **no loss of performance**
 - requires experience in both Julia and CUDA C programming (to understand error messages)

Why coding with Julia and CUDA.jl?

- Julia code for convolution kernel (zero padding, same size)

```
1 using CUDA ✓
2
3 function convolve(
4     source, destination,
5     filter, filter_norm::Int64,
6     dim1_s::Int64, dim2_s::Int64,
7     dim1_f::Int64, dim2_f::Int64
8 )
9     # Julia array is 1-based
10    i = (blockIdx().x-1) * blockDim().x + threadIdx().x
11    j = (blockIdx().y-1) * blockDim().y + threadIdx().y
12    if i <= dim1_s && j <= dim2_s
13        x0 = dim1_f ÷ 2 + 1
14        y0 = dim2_f ÷ 2 + 1
15        p0 = max(1+x0-i,1)
16        p1 = min(dim1_s+x0-i,dim1_f)
17        q0 = max(1+y0-j,1)
18        q1 = min(dim2_s+y0-j,dim2_f)
19        s = 0
20        for p = p0:p1
21            for q = q0:q1
22                @inbounds s += source[i+p-x0,j+q-y0] * filter[p,q]
23            end
24        end
25        destination[i,j] = (s ÷ filter_norm)
26    end
27    return
28 end
```



A diagram illustrating the execution flow of the code. A purple arrow points from the `gaussian_filter` function definition to its usage in the `convolve` function. A green arrow points from the `convolve` function to the `@time @cuda` macro. Two blue arrows point from the `CUDA.HostKernel{typeof(convolve...)}` expression back to the `convolve` function, indicating the kernel's role in the CUDA execution.

```
32 gaussian_filter(n,σ) = [floor(255*exp(-(i^2+j^2)/(σ^2))) for i=-(n÷2):(n÷2),j=-(n÷2):(n÷2)]
33
34 dims = (8000,4000) | (8000, 4000)
35 nGF = 31 | 31
36 GF = CuArray( gaussian_filter(nGF, 8.0) .|> Int64 ) | 31×31 CuArray{Int64, 2, CUDA.Mem.DeviceBuffer}
37 sumGF = sum(Array(GF)) | 50195
38
39 src = rand(0:255, dims) .|> Int64 | 8000×4000 Matrix{Int64}:
40 dst = zeros(Int64, dims) | 8000×4000 Matrix{Int64}:
41 cu_src = CuArray(src) | 8000×4000 CuArray{Int64, 2, CUDA.Mem.DeviceBuffer}:
42 cu_dst = CuArray(dst) | 8000×4000 CuArray{Int64, 2, CUDA.Mem.DeviceBuffer}:
43
44 t = 32 | 32
45 @time @cuda threads=(t,t) blocks=(dims[1]÷t+1,dims[2]÷t+1) convolve(
46     cu_src, cu_dst, GF, sumGF, dims[1], dims[2], nGF, nGF) | CUDA.HostKernel{typeof(convolve...)}
47
```

OUTPUT PROBLEMS DEBUG CONSOLE TERMINAL JUPYTER

0.000066 seconds (45 allocations: 2.234 KiB)
CUDA.HostKernel{typeof(convolve), Tuple{CuDeviceMatrix{Int64, 1}, CuDeviceMatrix{Int64, 1}, CuDeviceMatrix{Int64, Int64, Int64}}}(convolve, CuFunction{Ptr{Nothing}} @0x00000000006462890, CuModule{Ptr{Nothing}}

Why coding with Julia and CUDA.jl?

- Julia code to call cuBLAS / cuSolver routines
 - Example 1
 - CPU code : $\mathbf{A} * \mathbf{Y}$
 - GPU code : $\mathbf{cuA} * \mathbf{cuY} \mid\triangleright \mathbf{Array}$
 - Example 2
 - CPU code : $\mathbf{A} \setminus \mathbf{b}$
 - GPU code : $\mathbf{CuArray(A)} \setminus \mathbf{CuArray(b)} \mid\triangleright \mathbf{Array}$

RigorousCoupledWaveAnalysis.jl

- Enhanced Transmission Matrix (ETM) algorithm by Moharam [Moharam1995]
 - Backward and forward iteration
 - Numerical stability (key contribution of the [Moharam1995] paper)
- RigorousCoupledWaveAnalysis.jl by Jon Schlipf
 - GitHub repo: <https://github.com/jonschlipf/RigorousCoupledWaveAnalysis.jl>
 - Code is in good quality
 - Bottlenecks (identified with Julia profiling macro @profview)
 - #1: `eigen(M)` with **non-Hermitian M**
 - ==> **cuSolver cannot help :- ([need improvement on the algorithm level]**
 - #2: `A \ b` via Lapack (solution of the linear equation $A.x = b$)
 - ==> **Array(CuArray(A) \ CuArray(b)) via cuSolver** [verified, 10x~100x speed up]

RigorousCoupledWaveAnalysis.jl

- Enable GPU acceleration with CUDA.jl
 - Install CUDA.jl into the package project : <https://docs.julialang.org/en/v1/stdlib/Pkg/>
 - Minimal change leads to at least 10x performance gain for bottleneck #2

```
42 # 2. [forward iteration]
43 # compute the reflected wave and the forward wave in the first layer
44 # bottleneck
45 ψref,ψm1 = slicehalf( -cat([I;sup.V],F(em[1])*[em[1].X*(a[1]/b[1])*em[1].X;I],dims=2) \ ([I;-sup.V]*ψin) )
```

```
14 # 2. [forward iteration]
15 # compute the reflected wave and the forward wave in the first layer
16 # CUDA version : Array(CuArray(A) \ CuArray(b))
17 ψref,ψm1 = slicehalf(
18     Array(CuArray(-cat([I;sup.V],F(em[1])*[em[1].X*(a[1]/b[1])*em[1].X;I],dims=2)) \ CuArray([I;-sup.V]*ψin)))
19 )
```

```
etm_reftra_fast():
8.315675 seconds (334 allocations: 394.104 MiB, 0.19% gc time)
2.548487 seconds (30.57 k allocations: 1.263 GiB, 1.41% gc time)
---
```

```
etm_reftra():
8.519581 seconds (345 allocations: 394.104 MiB, 0.33% gc time)
26.735786 seconds (30.38 k allocations: 1.879 GiB, 0.90% gc time)
---
```

julia>

```
etm_reftra_fast():
1.641657 seconds (323 allocations: 85.186 MiB)
0.396935 seconds (14.41 k allocations: 277.851 MiB, 9.70% gc time)
---
```

```
etm_reftra():
1.541210 seconds (323 allocations: 85.186 MiB)
1.585677 seconds (14.24 k allocations: 414.954 MiB, 0.40% gc time)
---
```

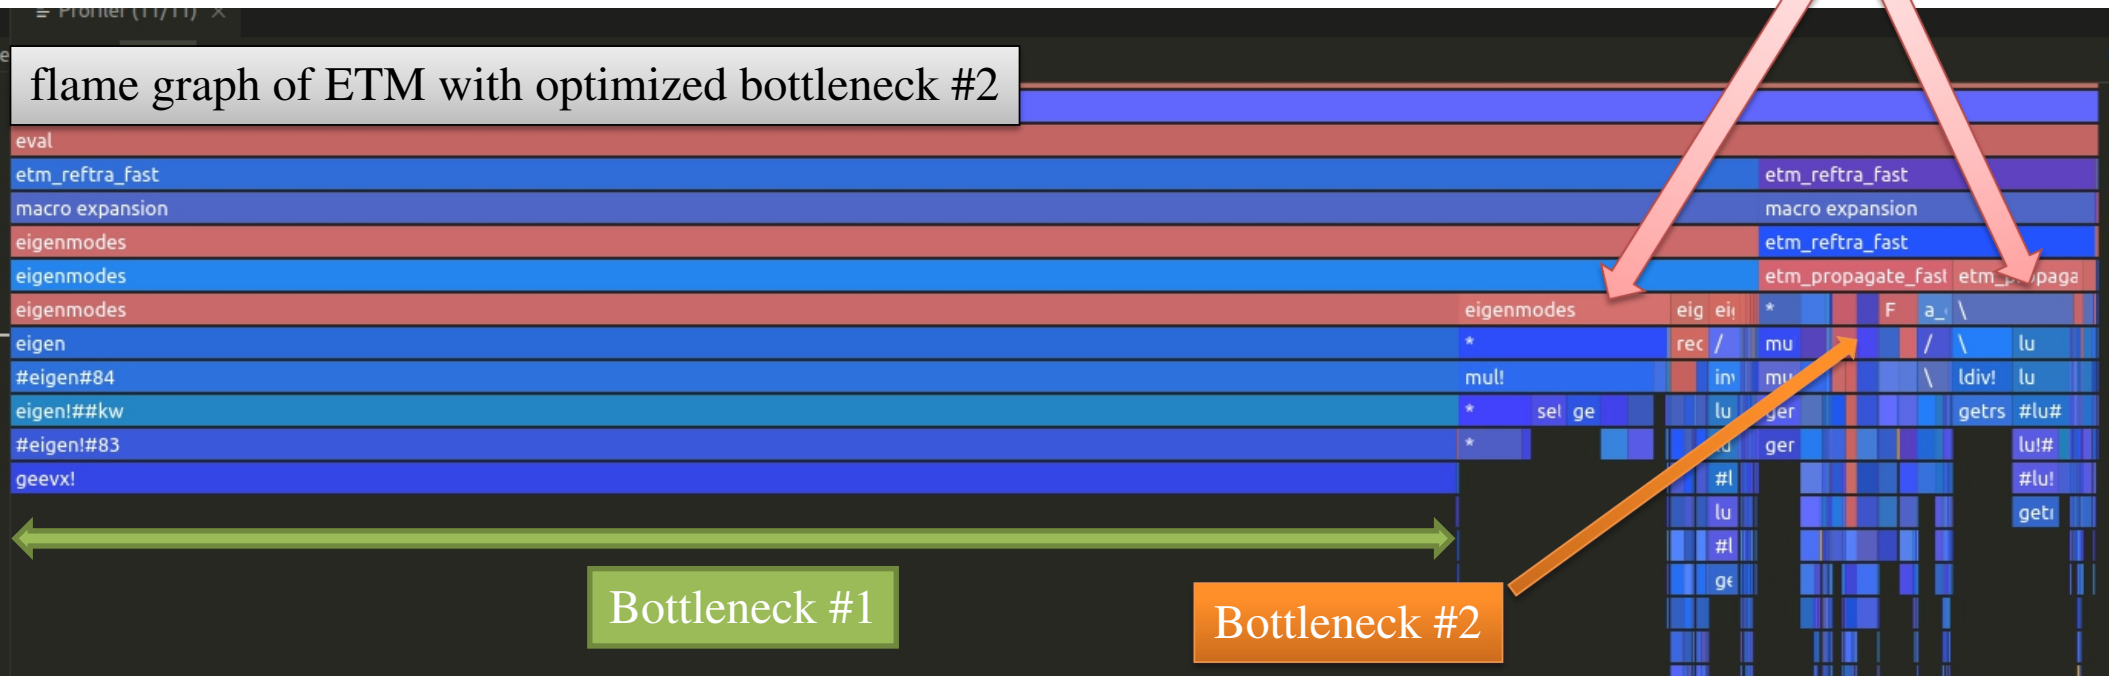
julia>

RigorousCoupledWaveAnalysisCUDA.jl

- Full implementation of the Enhanced Transmission Matrix algorithm with CUDA.jl
 - avoid memory transfer to achieve additional performance gain
 - excellent code design + strength of the polymorphism feature of Julia

Benefits from
full implementation

flame graph of ETM with optimized bottleneck #2

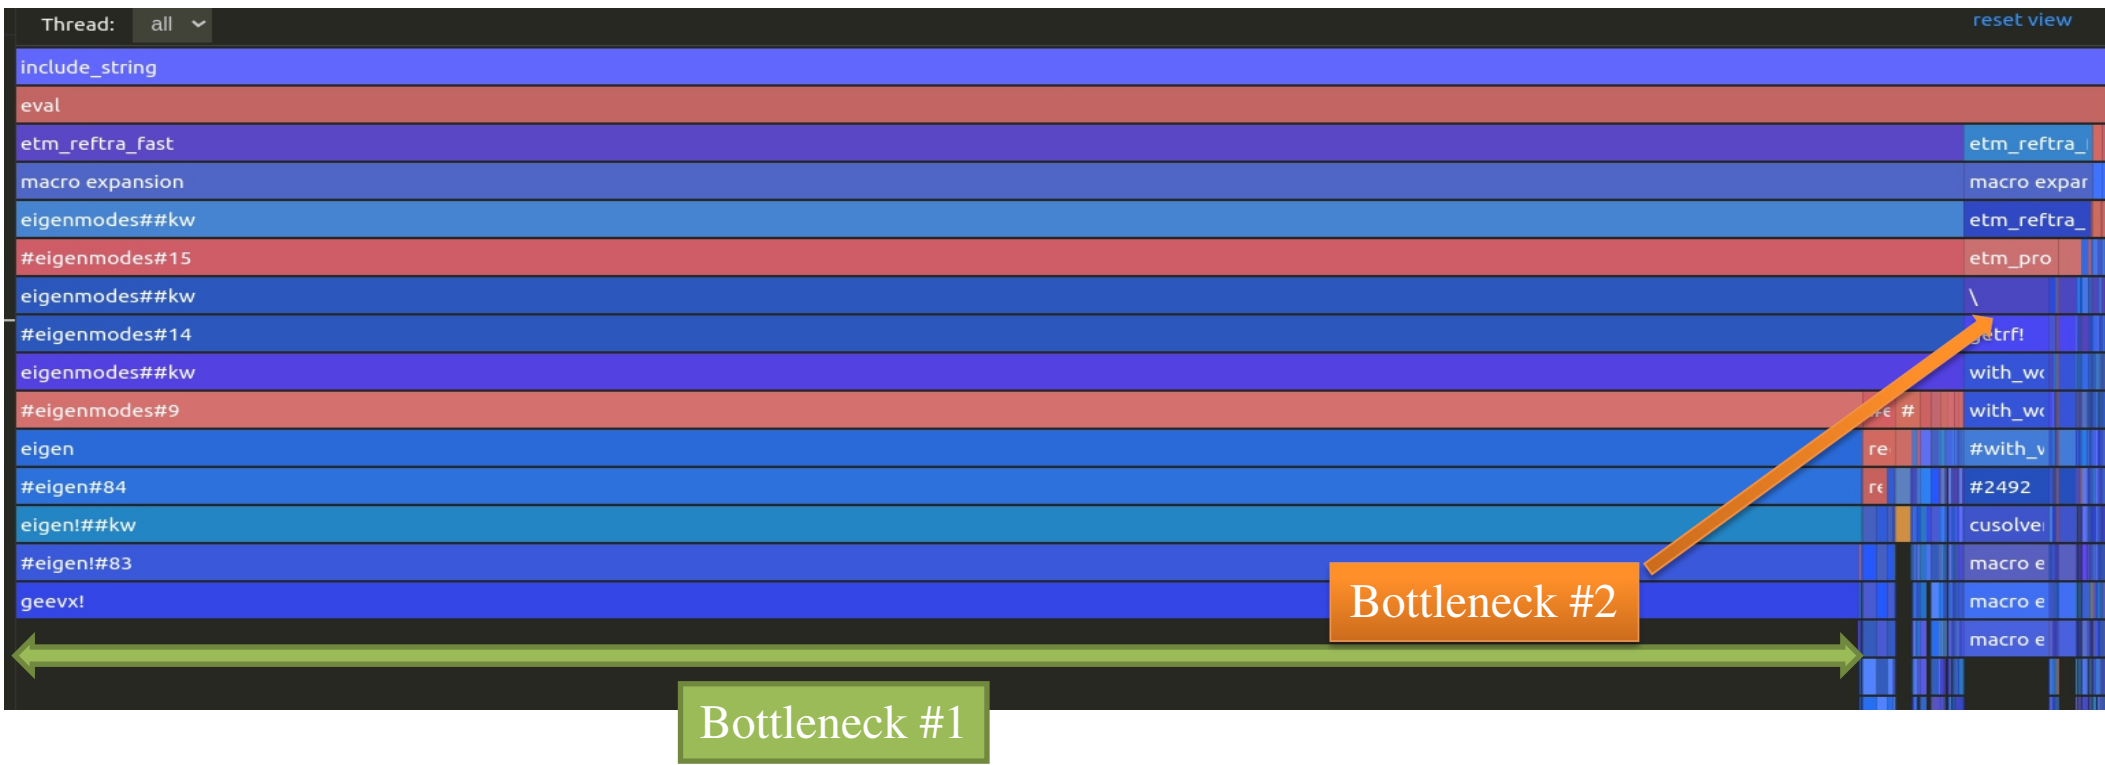


Bottleneck #1

Bottleneck #2

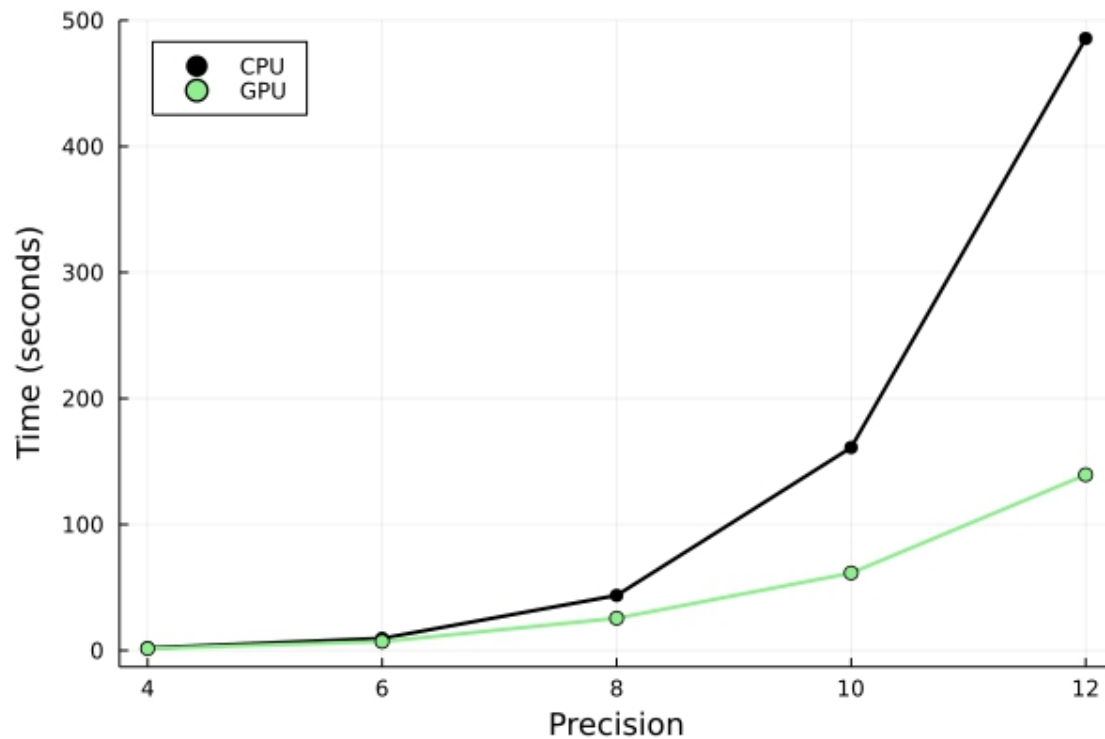
RigorousCoupledWaveAnalysisCUDA.jl

- Full implementation of the Enhanced Transmission Matrix algorithm with CUDA.jl
- Result of optimization: Bottleneck#1 now takes 85% execution time (60% before)



Test

- Test for correctness: against CPU code
 - integrated in the package test/runtests.jl
- Overall performance gain:
- *nvprof



References / Useful materials

FEM

FDTD

RCWA

Summary

- Common algorithms in computational electromagnetics often faces unique technical challenges
- Parallelization with GPU can improve, but has limitations such as memory size and inter-node communication
- Parallelized RCWA algorithm can easily be implemented in Julia programming language, with the help of CUDA.jl package
- The performance gain is significant despite that only one of the two bottlenecks is resolved.